

Context

The main factor in the success of MCMC methods is that they can be implemented with little efforts in a large variety of settings. Many softwares have been developed such as *WinBUGS* and *JAGS*, that helped to popularize Bayesian methods. These softwares allow the user to define his statistical model in a so-called BUGS language, then runs MCMC algorithms as a black box.

Although SMC methods have become a very popular class of numerical methods over the last 20 years, there is no such “black box software” for this class of methods. The *BiiPS* software aims at bridging this gap. From a graphical model defined in BUGS language, it automatically implements SMC algorithms and provides summaries of the posterior distributions.

SMC/particle methods

- Based on interacting particles systems governed by two stochastic mechanisms:
 - Mutation/Importance sampling: particles explore the space randomly and independently
 - Selection/Resampling: the best suited particles are duplicated, others removed
- Designed to sample from a sequence of distributions $\pi_k(x_{1:k}) = p(x_{1:k}|y_{1:k})$ when we can only compute the unnormalized version $\gamma_k(x_{1:k})$

$$\pi_k(x_{1:k}) = \frac{p(x_{1:k}, y_{1:k})}{p(y_{1:k})} = \frac{\gamma_k(x_{1:k})}{Z_k}$$

Generic SMC algorithm with N particles

- At time 1: for $i = 1, \dots, N$
 - Sample $x_1^{(i)} \sim q_1(x_1)$
 - Compute unnormalized weights $w_1^{(i)} = \frac{\gamma_1(x_1^{(i)})}{q_1(x_1^{(i)})}$ and estimate marginal likelihood $\hat{Z}_1 = \sum_{i=1}^N w_1^{(i)}$
- At time $k = 2, \dots, T$: for $i = 1, \dots, N$
 - Resample $\{x_{k-1}^{(i)}, w_{k-1}^{(i)}\}$ and set $w_{k-1}^{(i)} = 1/N$
 - Sample $x_k^{(i)} \sim q_k(x_k|x_{1:k-1})$
 - Compute unnormalized weights $w_k^{(i)} = w_{k-1}^{(i)} \frac{\gamma_k(x_{1:k}^{(i)})}{q_k(x_k^{(i)}|x_{1:k-1}^{(i)})}$ and estimate marginal likelihood $\hat{Z}_k = \hat{Z}_{k-1} \sum_{i=1}^N w_k^{(i)}$




Figure: Particles genealogical tree

Software features

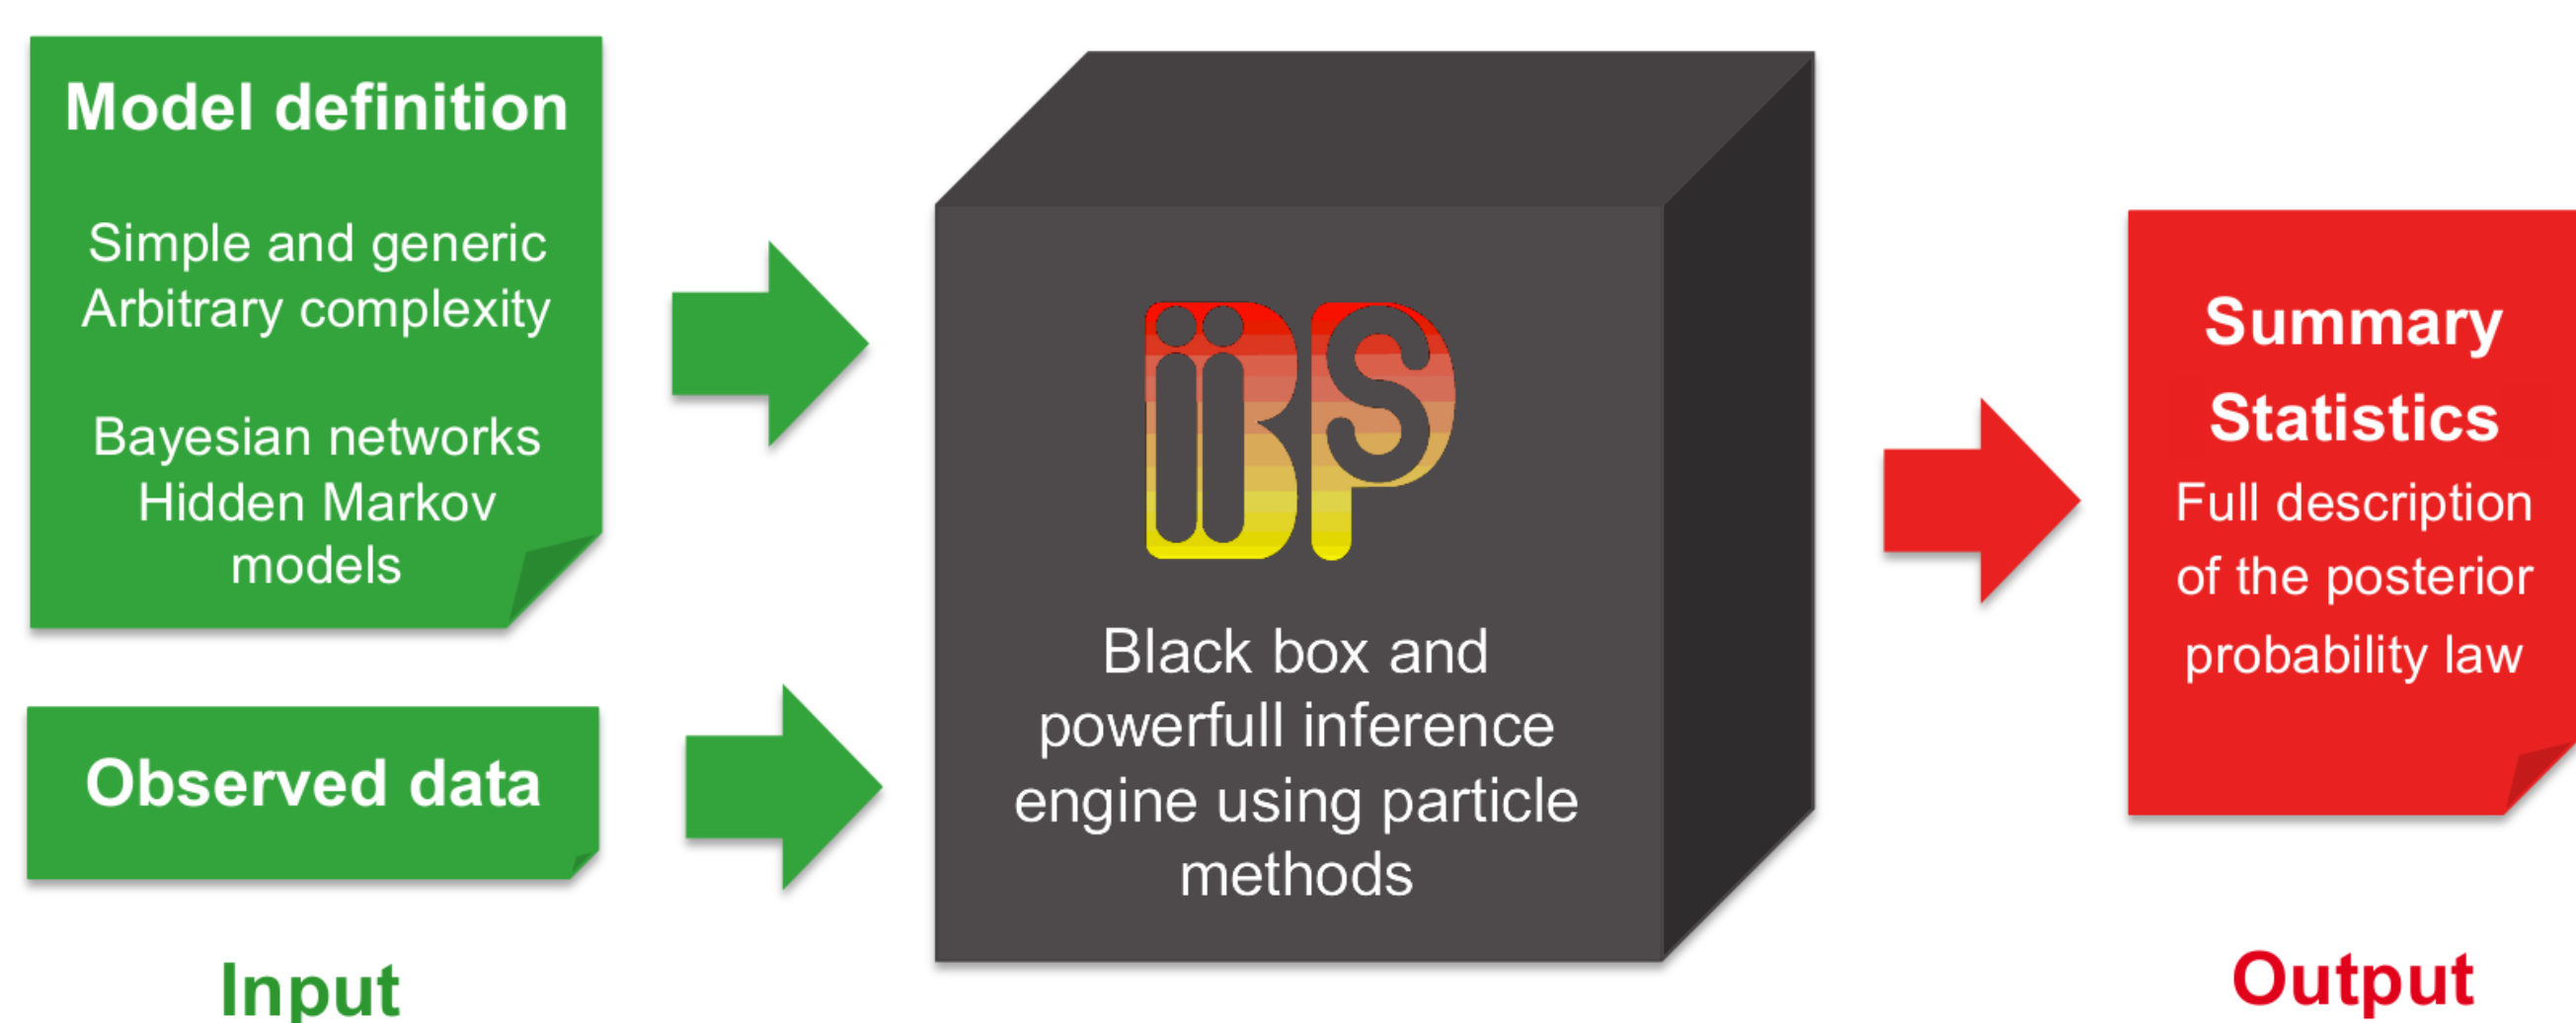


Figure: *BiiPS* input/output flowchart

BUGS language compatible

- Includes most usual uni/multivariate continuous/discrete distributions
- Standard operators, usual functions, matrix operations, ...

Development

- Free software adapted from *JAGS* © M. Plummer
- Core in C++ making use of Boost libraries
- R interface using *Rcpp* package
- Multi-platform: Linux, Windows, Mac

SMC techniques

- Forward filtering
- Backward smoothing
- Usual resampling algorithms: multinomial, residual, stratified and systematic
- Conditional sampler for Gaussian conjugate prior

Particle MCMC techniques

- Particle Independent Metropolis Hastings
- Particle Marginal Metropolis Hastings

Example in financial econometrics

Consider inferring the underlying volatility $x_{1:t}$ from observed price or rate data $y_{1:t}$

$$\begin{aligned} x_1 &\sim \mathcal{N}(0, \frac{\sigma^2}{1-\alpha^2}) \\ x_t &\sim \mathcal{N}(\alpha x_{t-1}, \frac{\sigma^2}{1-\alpha^2}) \quad t > 1 \\ y_t &\sim \mathcal{N}(0, \beta^2 \exp(x_t)) \quad t > 1 \end{aligned}$$

BUGS language

```
alpha ~ dunif(0, 0.99)
prec.x <- (1-alpha^2) / sigma^2
x[1] ~ dnorm(0, prec.x)
for (t in 2:t.max) {
  f[t] <- alpha * x[t-1]
  x[t] ~ dnorm(f[t], prec.x)
  prec.y[t] <- 1 / (beta^2 * exp(x[t]))
  y[t] ~ dnorm(0, prec.y[t])
}
```




Figure: Hidden Markov Model




Figure: Volatility simulation

RBiips package

Inference of the volatility

```
data <- list(t.max=100, sigma=1.0,
             alpha=0.91, beta=0.5,
             y=y)

# Compile the model and load the data
model <- biips.model("volatility.bug",
                    data)

# Run SMC algorithm
out.smc <- smc.samples(model, "x",
                       n.part=1000)

# Summary statistics
x.summ <- summary(out.smc$x,
                  fun=c("mean", "quantiles"),
                  probs=c(.05, .95))

plot(x.summ)

# Kernel density estimates
plot(density(out.smc$x, adjust=2))
```

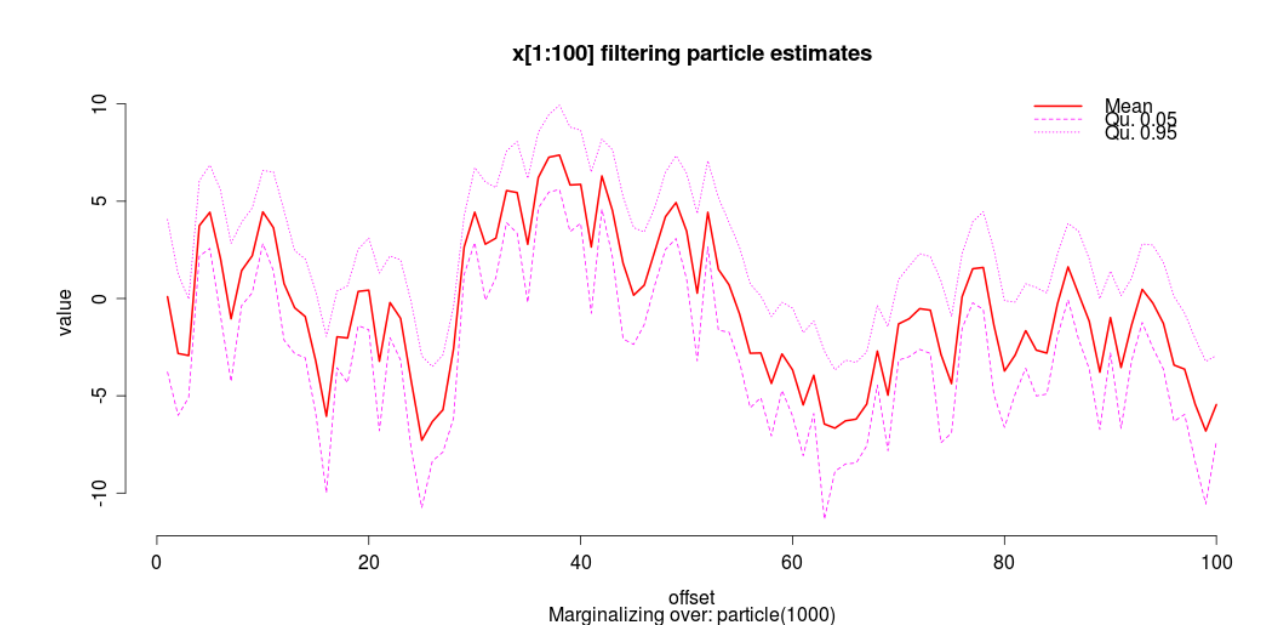


Figure: Summary statistics

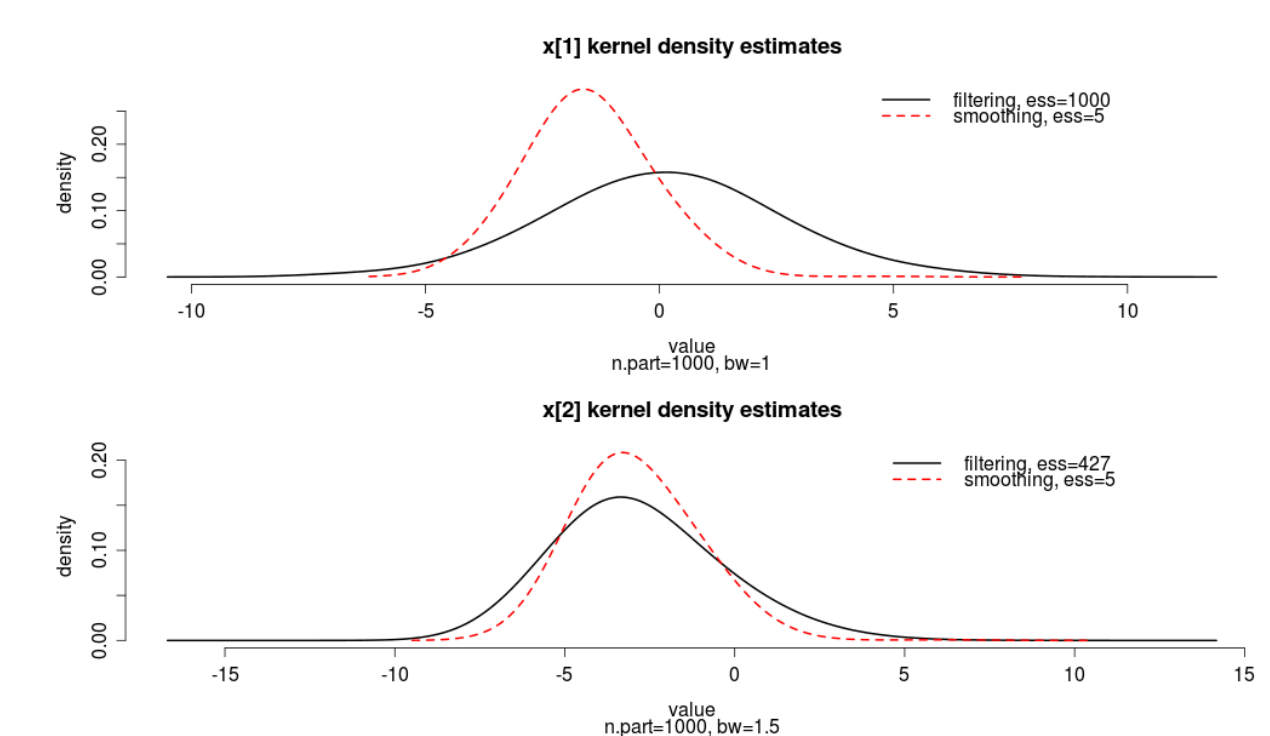


Figure: Kernel density estimates

Estimation of the fixed parameter α

```
data <- list(t.max=100, sigma=1.0,
             beta=0.5, y=y)

model <- biips.model("volatility.bug",
                    data)

# Sensitivity analysis
out.sens <- smc.sensitivity(model,
                           list(alpha=seq(0, .99, .01),
                           n.part=100))

plot(param$alpha,
     out.sens$log.marg.like)

# Burn in PMMH algorithm
update.pmmh(model, "alpha",
            n.iter=1000, n.part=100)

# Generate PMMH samples
out.pmmh <- pmmh.samples(model,
                          "alpha", n.iter=10000,
                          n.part=100)

# PMMH mean value
print(mean(out.pmmh$alpha))

# PMMH trace plot and histogram
plot(out.pmmh$alpha)
hist(out.pmmh$alpha)
```




Figure: α sensitivity analysis




Figure: α PMMH samples: trace plot and histogram

Future work

- Improve performance, parallelization, reduce memory footprint
- Interfaces: Matlab, Python, standalone
- More conjugate samplers, distributions and functions
- More advanced particle techniques
- Allow external user defined functions and samplers

References

- Andrieu, C., Doucet, A., and Holenstein, R. (2010). Particle markov chain monte carlo methods. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 72(3):269–342.
- Del Moral, P. (2004). *Feynman-Kac formulae: genealogical and interacting particle systems with applications*. Springer Verlag.
- Doucet, A., De Freitas, N., and Gordon, N. (2001). *Sequential Monte Carlo methods in practice*. Springer Verlag.
- Doucet, A. and Johansen, A. (2009). A tutorial on particle filtering and smoothing: Fifteen years later. *Handbook of Nonlinear Filtering*, pages 656–704.